PRE-SEED · NOVEMBER 2026 · INDEX0 RESEARCH LABS

INDEX0

Nothing merges until it runs.

status: qa_verified

3 passed / 0 failed

exit 0

sha256 9f2c…

--network none

approved

test_result

iteration 1

1.6 s

The bouncer for AI-written code.

github.com/INDEX0-RESEARCH-LABS/INDEX0

Writing got fast. Checking did not.

Three numbers explain the whole market.

90%

of developers now use AI at work.

DORA 2025, 5,000 respondents.

45%

of AI-written code ships a **security flaw**.

Veracode 2025, 100+ models.

89%

of a pull request's life is **waiting for review**.

CodePulse, 117,413 PRs.

So the writing got faster. The checking is still one person reading.

Sources: dora.dev/research/2025 · veracode.com 2025 GenAI Code Security Report · codepulsehq.com/research/cycle-time-decoded

So we built the door.

INDEX0 runs your tests, then says yes or no. Nothing else.

Ship or stop

One verdict per patch. **Not another report.**

1.6 seconds

The real suite, on a real repo. **Already measured.**

Sealed box

Code runs with **no network**. Nothing leaves.

Your keys

Your model contract. **Your bill.**

`services/agent-orchestrator · services/sandbox-manager · infra/litellm`

Three rules. That is the product.

No dashboard. No migration. No new CI to buy.

```
$ index0 verify

repository  acme/billing
suite       npm test
sandbox     firecracker

▸ architect  plan
▸ developer  patch
▸ critic     request_changes
▸ developer  patch
▸ critic     approve
▸ qa         3 passed / 0 failed

VERIFIED  exit 0 · 1.6 s · signed
```

Any repo that has a test command.

One command in the terminal you already use.

Exit 0, or the patch does not merge.

Read-only source mount · --network none · exit code 0 is the only pass

Six things a copycat cannot fake.

Every one came from a bug we shipped and then closed.

Runs the code

not just the diff

Seven seconds

warm cycle, end to end

Cannot cheat

test edits are dropped

No network

sealed sandbox

Your keys

your model, your bill

Signed proof

every run, hash-chained

The cheat guard exists because a Gemini run edited the very test that judged it.

What you can run today.

Actual screenshots from the repository at HEAD.

01 · THE COMMAND

```
$ index0 verify
repo  acme/billing
suite npm test

qa    3 passed
exit  0
time  1.6 s
```

One command in any terminal.

02 · THE RECEIPT

```
{
  "status": "qa_verified",
  "iterations": 1,
  "tests": {
    "passed": 3,
    "failed": 0
  },
  "exit_code": 0
}
```

One file anyone can re-run.

03 · THE PATCH

```
--- a/src/math.js
+++ b/src/math.js
 export function
   divide(n, d) {
+  if (d === 0) {
+    throw new TypeError
+  }
   return n / d;
 }
```

One diff you can actually read.

15 of 36 claimed capabilities are verifiable today — PROOF_STATUS.md, in the repo.

Where it pays for itself.

Four patches that all looked correct.



Security fixes

Compiles, passes review, **still ships a CVE**.



Big refactors

400 files where **nobody can read it all**.



Dependency bumps

Breaks a call site **you never opened**.



AI patches

Plausible, **not correct**.

Veracode 2025 GenAI Code Security Report · Stack Overflow Developer Survey 2025

A human reads 400 lines an hour.

The gate runs the whole suite in **1.6 seconds**. Same tests, same repo.

TODAY	WITH THE GATE
A reviewer reads **400 lines an hour**	**1.6 seconds** for the whole suite.
A cloud scanner means **your code leaves**	**Zero egress**: it runs where it lives.
A bug found in production costs **30× more**	**Caught before merge**, every time.
A review comment is **gone in a week**	**A signed receipt** that replays offline.

400 lines/hour: Cisco/SmartBear peer-review study · 30×: NIST 2002 · 1.6 s measured in demo/workspace, run 20261031T041126Z

Boring on purpose.

No research breakthrough required.



Runs anywhere

Docker or Firecracker. Your cloud or your laptop.



Any model

Swap vendors without touching the gate.

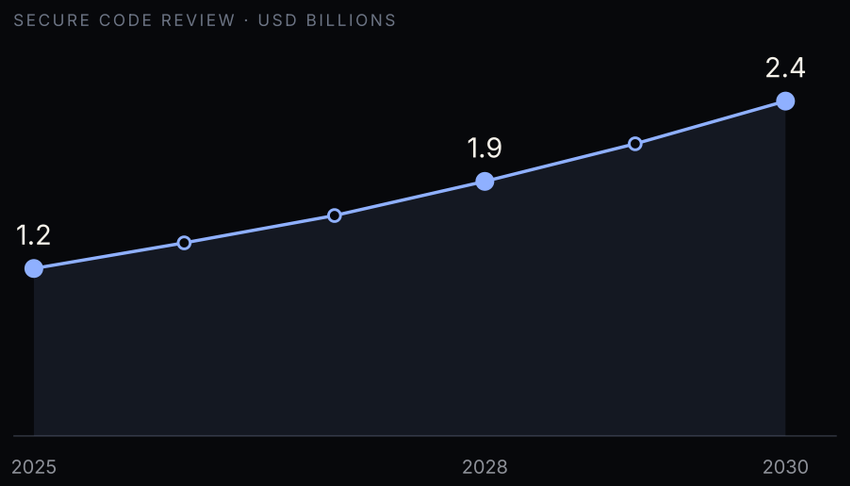


No patent

The public audit is the moat. Anyone can check it.

The budget is real. Here is the math.

Two published points. The four in between use the same 14.9% rate.



2,500	regulated organisations
× 500	developers each
× $360	per developer, per year

$450M we can sell into

That is 37% of the $1.22B category — the regulated half. We quote the smaller number.

Mordor Intelligence, Secure Code Review Platforms Market: USD 1.22B (2025) to USD 2.44B (2030), 14.88% CAGR. Page updated Sep 2026.

Three endings. All possible.

Every figure is accounts × $180,000, against $450M we can sell into.

Best case

The default gate. **1,500 accounts**, 60% of the segment.

$270M ARR

Realistic

5% of 2,500 organisations. **120 accounts**.

$21M ARR

Worst case

A 10-person shop, one sector. **15 accounts**.

$2.7M ARR

No revenue today. These are outcomes of the price, not a forecast.

Two ways in.

We do not buy attention. We hand over proof.



Three design partners

Platform and AppSec leads at 200–5,000-developer firms. **Founder-led**, three at a time.



The receipt is the ad

Every run leaves a signed file. **Proof travels**; marketing does not.

A receipt names the repo, the suite, the exit code and the hash. It gets pasted into tickets.

Four months. Four proofs.

Each one is in the repository, with the commit that produced it.

OCT 2026

License audit clean. No copyleft in the paid tree.

OCT 2026

30 of 30 tests green. Four tiers, one loop.

31 OCT 2026

Six verified runs. Three on a free model route.

NOV 2026

Price set. Demo and proof pack shipped.