PRE-SEED · NOVEMBER 2026 · INDEX0 AI INC.

INDEX0 AI

Own your AI engineer.

surfaces: cli · ide · web

status: qa_verified

3 passed / 0 failed

exit 0

--network none

self-hosted

your model key

sha256 9f2c…

1.6 s

People Over Tools. Work Verified. Time to Unplug.

github.com/INDEX0-RESEARCH-LABS/INDEX0

AI writes the code now. Nobody checks it.

Three numbers explain the whole market.

90%

of developers now use AI at work.

DORA 2025, 5,000 respondents.

45%

of AI-written code ships a **security flaw**.

Veracode 2025, 100+ models.

89%

of a pull request's life is **waiting for a human**.

CodePulse, 117,413 PRs.

The generator arrived. The proof did not.

Sources: dora.dev/research/2025 · veracode.com 2025 GenAI Code Security Report · codepulsehq.com/research/cycle-time-decoded

One platform. Three surfaces. One rule.

INDEX0 AI writes the code where you already work, then proves it runs.

Terminal

The **index0** CLI: plan, review, run agents, from where you type.

Editor

A VS Code panel for runs, diffs and **inline completions**.

Browser

A workbench and portal. **No install**, nothing to set up.

And one rule

Nothing merges until **the tests run**.

Four agents. One verdict.

Ask in any surface. The same loop answers, and the last tier runs the code.

```
$ index0 run "reject divide by zero"

surface     cli
repository  acme/billing
suite       npm test
sandbox     firecracker

▸ architect   plan
▸ developer   patch
▸ critic      request_changes
▸ developer   patch
▸ critic      approve
▸ qa          3 passed / 0 failed

VERIFIED  exit 0 · 1.6 s · signed
```

Any repo that has a test command.

Any surface: terminal, editor, browser.

Exit 0, or the patch does not merge.

Read-only source mount · --network none · exit code 0 is the only pass

Six things a copycat cannot fake.

Every one came from a bug we shipped and then closed.



Runs the code

not just the diff



Self-hosted

your cloud or your laptop



Cannot cheat

test files hashed, then compared



No network

sealed sandbox



Your keys

your model, your bill



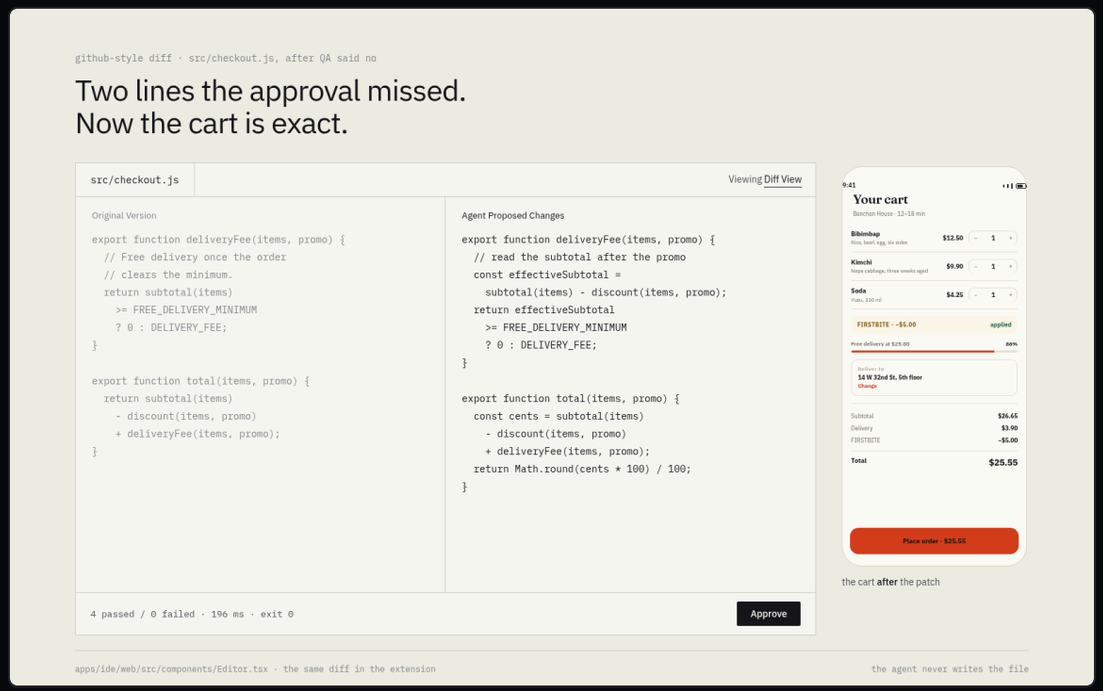
Signed proof

every run, hash-chained

Test, evaluator and harness files are hashed before a run; a diff marks it INVALID_RUN.

One cart. One run. One refusal.

A food-delivery checkout, a $5 promo, and a fee that still read the pre-promo cart.



04 · THE PATCH

Two lines: read the payable subtotal, round the total. 4 passed, 196 ms.

Terminal. The critic approved the diff. QA exited 1.

`3 passed / 1 failed · 2530 ms`

Editor. Both versions on screen. You approve.

`Original Version · Agent Proposed Changes`

Browser. A plan in English, the failure in the open.

`Execution Plan · INDEX0 Agent & Sandbox Terminal`

`Design mockups from the shipped components, not screen captures · all six frames in pitch/shots.html`

`37 CLI tests · 94 web-workbench tests · 6 extension tests · 16 of 37 capabilities verifiable (PROOF_STATUS.md)`

Where it pays for itself.

Four patches that all looked correct.



Security fixes

Compiles, passes review, **still ships a CVE**.



Big refactors

400 files where **nobody can read it all**.



Dependency bumps

Breaks a call site **you never opened**.



AI patches

Plausible, **not correct**.

Veracode 2025 GenAI Code Security Report · Stack Overflow Developer Survey 2025

Their assistant suggests. Yours decides.

Same patch. One tool stops at the diff; the other runs your suite before it lands.

CLOUD CODING ASSISTANTS	INDEX0 AI
They complete code **inside their cloud**	**It runs in yours**: VPC, laptop, air-gap.
They hand you a diff and **stop there**	**QA runs your suite** before it lands.
They bill **per seat, every month**	**You own the deployment** and the model key.
Their review comment **fades in a week**	**A signed receipt** that replays offline.

1.6 s measured on demo/workspace · run 20261031T041126Z · 3 passed / 0 failed

Boring on purpose.

No research breakthrough required.

Runs anywhere

Docker, Firecracker or E2B. **Your cloud** or your laptop.

Any model

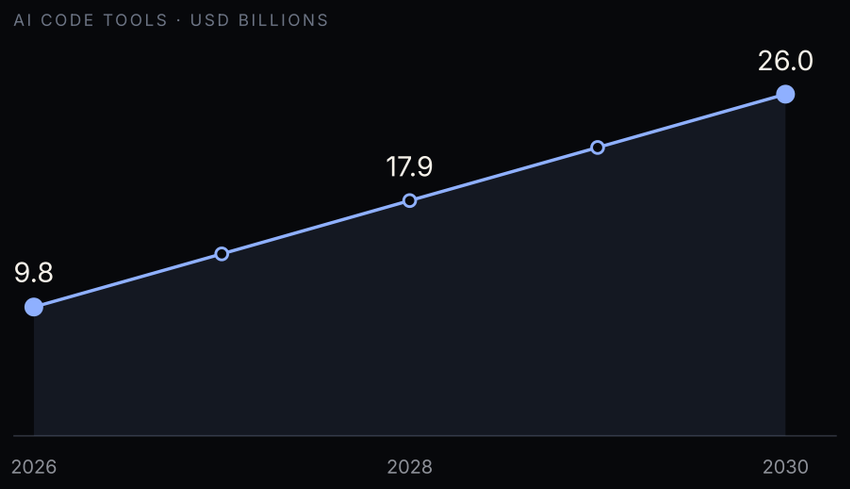
Swap vendors without touching the loop.

No patent

The public audit is the moat. Anyone can check it.

The budget is real. Here is the math.

AI code tools, 2026 to 2030. We sell the half that must run on-premises.



Published: 2026 and 2030 Interpolated between them

2,500	regulated organisations
× 500	developers each
× $240	per developer, per year

$300M we can sell into

$20 a seat a month, the price in our own model. That slice is 3% of the 2026 category.

Grand View Research, AI Code Tools Market: USD 9.8B (2026) to USD 26.0B (2030), 27.1% CAGR, updated Jun 2026. GVR names on-premises the fastest-growing deployment segment, driven by data-sovereignty rules.

Three endings. All possible.

Every figure is accounts × $120,000, against the $300M we can sell into.

Best case

The default gate in regulated engineering. **1,500 accounts**, 60% of the segment.

$180M ARR

Realistic

5% of 2,500 organisations. **120 accounts**.

$14.4M ARR

Worst case

A 10-person shop, one sector. **15 accounts**.

$1.8M ARR

No revenue today. These are outcomes of the price, not a forecast.

Two ways in.

We do not buy attention. We hand over proof.



Three design partners

Platform and AppSec leads at 200–5,000-developer firms. **Founder-led**, three at a time.



The receipt is the ad

Every run leaves a signed file. **Proof travels**; marketing does not.

A receipt names the repo, the suite, the exit code and the hash. It gets pasted into tickets.

Four months. Four proofs.

Each one is in the repository, with the commit that produced it.

OCT 2026

Three surfaces shipped. CLI, editor, browser. All wired to the same orchestrator.

OCT 2026

Review loop closed. 30 of 30 orchestrator tests. QA gate cannot fake a pass.

NOV 2026

Benchmark run. 20 tasks, hidden tests, same model. Baseline 89%, INDEX0 74%. Reported honestly.

NOV 2026

Claude via Bedrock. One-command demo. 33 s end-to-end. `sandbox: docker (exit 0)`.

Raising $750K · 18 months · three paid pilots · 16 of 37 capabilities verifiable · benchmarks/verified-bench/RESULTS.md