PRE-SEED · OCTOBER 2026

INDEX0

The verification layer for AI-written code — running in your tenant, on your keys, with proof you can replay.

A working product, a published audit of our own claims, and no benchmark we can't reproduce.

01 — THE PROBLEM

AI made code cheap. Knowing it's correct stayed expensive.

10×

more security findings from AI-assisted developers, who also wrote 3–4× more code.

45%

of AI-generated samples introduced an OWASP Top 10 flaw — and it did not improve across model generations.

86%

of AI-generated XSS samples failed. Java fared worst at 72% of samples flawed.

The bottleneck moved from **writing** code to **proving** it. Every team now reviews more code, faster, with the same number of senior eyes.

Apiiro, *4× velocity, 10× vulnerability*, 4 Sep 2025 · Veracode, *2025 GenAI Code Security Report*, 30 Jul 2025

02 — THE SYMPTOM

"Almost right" is the expensive failure mode.

46%

of developers actively distrust AI tool accuracy; 33% trust it.

66%

say the frustration is code that's "almost right, but not quite".

45%

say debugging AI-generated code takes longer than code they wrote.

Building it is no longer the hard part. **Accepting it is.**

Stack Overflow Developer Survey 2025, n≈31,500 · published Jul 2025

03 — THE MEASUREMENT

Speed went up. Stability went down.

−7.2%

delivery stability for a 25% increase in AI adoption.

−1.5%

delivery throughput at the same adoption level.

2025

throughput turned positive — instability is still negative.

DORA's own reading: AI amplifies whatever your delivery system already does. If verification is weak, AI makes it weaker, faster.

DORA / Google Cloud, *Impact of GenAI in Software Development*, 2024 · updated Sep 2025

04 — WHY NOW

Sovereignty stopped being a preference and became a procurement score.

JAN 2025 **DORA in force** for EU financial entities — ICT third-party risk, 4-hour incident reporting.

OCT 2025 **EU Cloud Sovereignty Framework v1.2.1** scores sovereignty in procurement, backed by a €180M tender.

NOV 2025 **India DPDP Rules notified** — Significant Data Fiduciaries can be ordered to keep data in India; penalties to ₹250 crore.

2025 → 2027 **Sovereign cloud IaaS** grows $59.3B → $110.6B globally; Europe $6.9B → $23.1B.

"Send your source code to a US SaaS" is a non-starter in exactly the accounts with the most to lose.

European Commission, 20 Oct 2025 · India MeitY, 13 Nov 2025 · Gartner, 9 Feb 2026

05 — THE MARKET

The budget is real. The estimates are soft, so we show the range.



AI coding tools, 26.2% CAGR. Two independent analyst houses differ by **~4×** — so we cite the range, not a point, and we never quote a TAM for "AI code verification" because none exists.

Our budget line is **AI governance**: AppSec and audit, not the IDE.

Mordor Intelligence, 2026 · MarketsandMarkets · Gartner: 90% of enterprise engineers on AI code assistants by 2028

06 — THE PRODUCT

A gate, not another generator.

TIER 1

Architect

Turns the objective into invariants that can be violated.

TIER 2

Developer

Edits the real files. No hunk arithmetic, no stubs.

TIER 3

Critic

Adversarial review. Only critical/error findings can block.

TIER 4

QA

Runs the candidate code against the repo's suite in a network-isolated container.

EVIDENCE The exit code decides — not a model's opinion of the diff.

SIGNED Every cycle is recorded, hash-chained and attributable to a key.

YOURS Your cloud, your model keys, your source never leaves.

07 — THE DEMO

This ran this morning, on a free model.

```
$ ./demo/run-demo.sh
==> litellm gateway (free route, no database)
    model groq reachable
==> sandbox-manager (provider=docker, image=node:22)
==> agent-orchestrator
==> running the review cycle

model      : groq
status     : qa_verified
iterations : 1
critic     : approve
qa passed  : True
tests      : 3 passed / 0 failed in 1332 ms
sandbox    : live_microvm (exit 0)
history
  architect   plan_proposal
  developer   implementation
  critic      review_verdict   approve
  qa          test_result      passed
```

The patch only counts because QA **executed** it — no network, read-only mount, 1.3 seconds. This is the whole diff it shipped:

```
--- a/src/math.js
+++ b/src/math.js
@@ -1,4 +1,7 @@
 export function divide(numerator, denominator) {
+  if (denominator === 0) {
+    throw new RangeError('division by zero');
+  }
   return numerator / denominator;
 }
```

---

demo/runs/20261031T002004Z — request, cycle transcript and diff committed with the code

08 — THE BUYER

200–5,000 developers, regulated, already on Azure.

Who signs

PLATFORM Owns the engineering system and the AI rollout.

APPSEC Owns the finding backlog we're about to make worse.

CISO Owns the answer to "where did that code go?".

Why they can't buy it today

FINANCE DORA third-party risk. Inference must stay in-perimeter.

HEALTH Source and patient-adjacent logic under HIPAA.

PUBLIC FedRAMP High, IL5/IL6, ITAR, EU data residency.

Honest caveat: self-hosting is not whitespace — Qodo and OpenHands already ship it. The wedge is an enforced gate with replayable evidence at seat counts incumbents won't serve.

09 — COMPETITION

Everyone reviews. Nobody hands you the receipt.

TOOL	WHAT IT IS	VERIFICATION GAP	SELF-HOST
Cursor	AI IDE + Bugbot	Reviews with the model family that generated the code	No documented
GitHub Copilot	IDE + PR + CodeQL	Strong static scanning; silent on whether tests ran	Cloud inference
CodeRabbit	PR review bot	Comment-level; no merge-blocking artifact	Enterprise, 500+ seats
Qodo	Review + test gen	Review-centric; the gate isn't the product	**Yes, air-gapped**
OpenHands	MIT agent platform	Runs tests; no independent Architect/Critic/QA split	**Yes**
INDEX0	**Verification gate**	**Plan → diff → sandboxed test → verdict, signed per commit**	**Yes, your keys**

The gap nobody closes: a record you can **replay offline** and verify without trusting the vendor. That's the claim we're asking you to fund.

Pricing and roadmap claims deliberately omitted — we did not verify them. Nothing here is sourced from a vendor's marketing page.

10 — WHAT IS ACTUALLY VERIFIED

We audited our own claims and deleted the ones we couldn't reproduce.

15/36

claimed capabilities verifiable today, listed line by line in **PROOF_STATUS.md**, shipped in the repo.

0

published benchmark scores. The old 96.2% SWE-bench and 92% AiderBench numbers were unverified and removed.

30/30

orchestrator tests green, plus 14 sandbox-manager and the full monorepo build.

> "We are building a verification product. The first thing we verified was ourselves — and any vendor selling you trust should be able to show you the same page."

PROOF_STATUS.md · benchmarks/CLAIM_AUDIT.md · audited 28 Oct 2026

11 — THE ASK

Pre-seed: $750K, eighteen months, three paid pilots.

What the round buys

01 One paid pilot in a regulated customer’s own tenant.

02 Three design partners on the evidence format.

03 A replayable artifact accepted by a real audit.

What we will not claim

- No benchmark score until we can reproduce it in public.
- No revenue, no users, no logos on this deck.
- No margin until inference is metered on a real tenant.

$750K across 18 months: two hires, three design partners, SOC 2 Type I. Milestone is a signed, replayable artifact — not a leaderboard position.

APPENDIX

The five things that can kill us.

01 **Verification is a feature, not a company.** Cursor, GitHub and Qodo can bundle review-plus-tests into seats they already own.

02 **No forcing function.** "Audit-ready evidence" is a nice-to-have until an auditor demands it. If that demand is 2029, we don't reach it.

03 **Two people cannot pass enterprise procurement.** SOC 2, pen test, 9–18 month cycles — runway ends first.

04 **Credibility debt is real.** The audit found an OTP bypass, injectable telemetry and a review path that could return approvals offline. All fixed and pinned by tests — but being caught hiding them would have been fatal.

05 **Own-keys means no telemetry.** No usage data, no expansion signal, no ROI proof we can show a board.

Stated plainly, because you will find them anyway. Mitigations are in the notes; ask us about any one of them.